

# Week 9 Support Vector Machine (SVM)

Theory and Practice

Ying Lin, Ph.D  
March 13, 2020

# Overview

- Linearly separable data set and margin maximization
- Non-linearly separable data set soft margin classifier
- Soft margin classifier vs. hard margin classifier
- Kernel trick
- R/Python demo

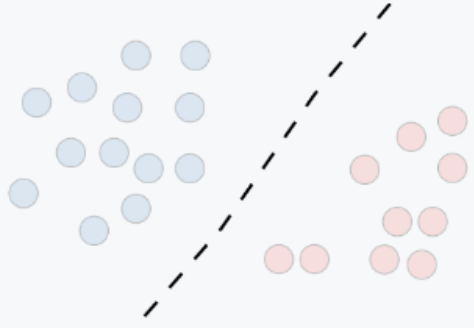
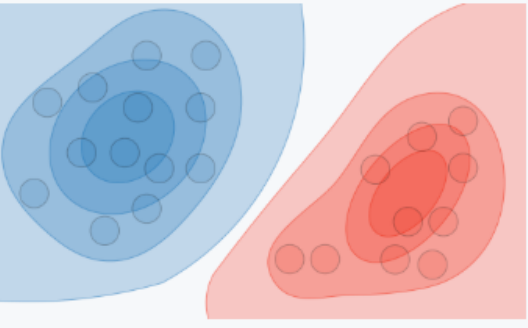
# Support Vector Machine (SVM) Motivations

- Work with both linearly separable and linearly non-separable cases
- Maximize margin between categories
- Margins are only decided by the closest data points (support vectors)
  - SVM is robust for the outliers
  - Distance between data points and decision boundary indicates confidence of prediction

# Generative vs. Discriminative Algorithms


- Generative algorithm: infer the distribution that generate the observed data (joint probability)
  - Model probability of individual classes
  - Example: naive Bayes classifier,
- Discriminative algorithm: identify the decision boundary that separates different classes of the observed data. Make fewer assumptions about the data distribution than generative algorithm
  - Model the (hard or soft) decision boundary between different classes
  - Example: Support Vector Machine, logistic regression, decision tree induction
- Discriminative algorithm usually outperforms generative algorithm, given large enough training dataset

# Generative vs. Discriminative Algorithms (2)

|                       | Discriminative model  | Generative model  |
|-----------------------|---|---|
| <b>Goal</b>           | Directly estimate $P(y x)$  | Estimate $P(x y)$ to then deduce $P(y x)$   |
| <b>What's learned</b> | Decision boundary   | Probability distributions of the data   |
| <b>Illustration</b>   |  |  |
| <b>Examples</b>       | Regressions, SVMs   | GDA, Naive Bayes  |


# SVM Hyperplane

- Linearly separable cases
- Maximize the margin between two parallel hyperplanes or minimize  $||\vec{w}||$
- Support vectors: data points on those two parallel hyperplanes
  - Number of support vectors: complexity of models
- Prediction: hyperplane divides the feature space in half and we can classify new points by determining which side of the hyperplane they fall on



# Soft Margin Classifier

- Allow SVM to make a certain number of mistakes ( $\xi$ )
- Keep margin as wide as possible so that rest of points can still be classified correctly
- $y_i(\vec{w} \cdot \vec{x}_i + b) \geq 1 - \xi_i$




# Loss Function for SVM: Hinge Loss


- SVM loss leads to SVM model with the max-margin property
- Hinge/SVM loss
  - $Hinge(y_i, \hat{y}_i) = \sum_i \max(0, 1 - y_i * \hat{y}_i)$
  - Scenario 1:  $y_i = -1, \hat{y}_i = 0.5$ , Hinge loss is 1.5
  - Scenario 2:  $y_i = -1, \hat{y}_i = -0.2$ , Hinge loss is 0.8
  - Scenario 3:  $y_i = -1, \hat{y}_i = -1.6$ , Hinge loss is 0
  - Therefore, Hinge loss penalizes those confident misclassification the most, and no penalty on the correct prediction



# Linearly Inseparable Challenge



# Linearly Inseparable Challenge in a Higher Dimension

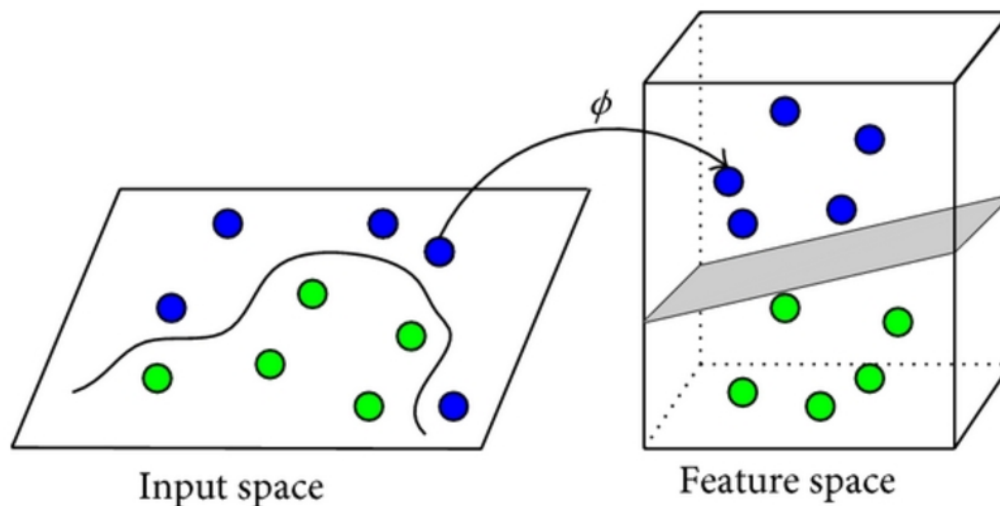


# Kernel Tricks

- Deal with classification that requires multiple (and nonlinear) boundaries
- Ideas: apply mathematical functions (kernels) to project linearly inseparable data points into higher dimensional space so that a hyperplane could be found to separate different classes
  - $K(X_i, X_j) = \phi(X_i) \cdot \phi(X_j)$
  - Why kernel function is a measure of similarity: the inner product above means the projection of  $\phi(X_i)$  on  $\phi(X_j)$  or how much overlap (or how similar) between  $\phi(X_i)$  and  $\phi(X_j)$
  - Kernel function is a distance function on a higher dimension where hyperplane could be found.
  - Project the data to the higher dimension space in order to find hyperplane to separate classes without explicitly projecting with  $\phi$  function and calculate the distance in the projected higher dimension space

# Kernel Trick (2)

- The kernel trick avoids the explicit mapping that is needed to get linear learning algorithms to learn a nonlinear function or decision boundary.
- For all  $x$  and  $y$  in the input space, certain functions can be expressed as an inner product in another space. The function is often referred to as a kernel or a kernel function.



# Non-linear Kernel Functions

- SVM is about finding a hyperplane (through margin maximization) to separate different classes
- Kernel: find distance between data points on a higher dimensional space where a hyperplane could be identified
  - Radial basis function (RBF): project to infinite dimensional space

$$K(X_m, X_n) = \exp\left(-\frac{||X_m - X_n||^2}{2\sigma^2}\right)$$

- Polynomial kernel

$$K(X_m, X_n) = (X_m \cdot X_n + c)^d$$

- Sigmoid kernel

$$K(X_m, X_n) = \tanh(\kappa X_m \cdot X_n - \delta)$$

- Compare above distance functions with Manhattan and Euclidean distance functions in the current feature space prior to projecting to the higher dimensional space

# Data Preprocessing and Preparation

- Kernel functions are based on distance (distance on higher dimensions) and herefore require the following two preprocessing techniques
- Normalization
- Conversion from categorical attributes to numerical attributes
- Missing value imputation

# Model Parameters

- $C$  (Cost): used for regularization, penalty associated with misclassification
  - Soft margin classification: allow a small number of misclassifications
  - Hyperparameter controlling penalizing use of slack variables  $\xi$
  - Higher  $C$ : lower bias but risk of overfitting; higher complexity of decision boundary; higher number of support vectors; Lower  $C$ : higher bias but lower variance
- Kernel function specific parameters
  - $\gamma$  for RBF: controls the tradeoff between error due to bias and variance
    - $k(X_m, X_n) = \exp(-\gamma ||X_m - X_n||^2)$
    - Higher  $\gamma$  means higher risk of overfitting; lower  $\gamma$  means higher bias; moves in the same direction as  $C$
  - Degree for Polynomial kernel
  - Scale for Polynomial kernel

# Properties of SVM

- Pros
  - Excellent performance for a wide variety of tasks because it can deal with nonlinear decision boundary
  - Provide confidence estimate (probability transformed from distance by proper mathematical functions)
  - Robust to noisy data
  - Flexible with data representation: numerical and categorical attributes
  - Effective with low sample/feature ratio data
- Cons
  - Lack of interpretability. Black-box model
  - Risk of choosing wrong kernel function



# Demo Dataset: Diabetes

```
# install.packages("mlbench")
library(caret)
library(mlbench)
data("PimaIndiansDiabetes")
diabetes <- PimaIndiansDiabetes
str(diabetes)
```

```
## 'data.frame':    768 obs. of  9 variables:
##  $ pregnant: num  6 1 8 1 0 5 3 10 2 8 ...
##  $ glucose : num  148 85 183 89 137 116 78 115 197 125 ...
##  $ pressure: num  72 66 64 66 40 74 50 0 70 96 ...
##  $ triceps : num  35 29 0 23 35 0 32 0 45 0 ...
##  $ insulin : num  0 0 0 94 168 0 88 0 543 0 ...
##  $ mass     : num  33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
##  $ pedigree: num  0.627 0.351 0.672 0.167 2.288 ...
##  $ age      : num  50 31 32 21 33 30 26 29 53 54 ...
##  $ diabetes: Factor w/ 2 levels "neg","pos": 2 1 2 1 2 1 2 1 2 2 ...
```

```
summary(diabetes)
```

```
##      pregnant      glucose      pressure      triceps
## Min.      : 0.000   Min.      : 0.0   Min.      : 0.00   Min.      : 0.00
```

17/35

# Train and Test Data Partition

```
library(caret)
```

```
set.seed(188)
```

```
train_index <- createDataPartition(diabetes$diabetes, p = 0.7, list = FALSE)
```

```
diabetes_train <- diabetes[train_index, ]
```

```
diabetes_test <- diabetes[-train_index, ]
```

# Demo: Train a Linear SVM Model

```
set.seed(1818)
model_svm_linear <- train(diabetes ~ ., data = diabetes_train,
                          method = "svmLinear",
                          preProcess = c("center", "scale"),
                          trControl = trainControl(method = "boot", number = 25),
                          tuneGrid = expand.grid(C = seq(0, 1, 0.05)))
```


```
model_svm_linear
```

```
## Support Vector Machines with Linear Kernel
##
## 538 samples
## 8 predictor
## 2 classes: 'neg', 'pos'
##
## Pre-processing: centered (8), scaled (8)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 538, 538, 538, 538, 538, 538, ...
## Resampling results across tuning parameters:
##
##  C      Accuracy  Kappa
##  0.00      NaN      NaN
##  0.05  0.7617913  0.4563736
##  0.10  0.7600813  0.4535191
```

19/35

# Sensitivity Test for C

```
plot(model_svm_linear)
```



# Linear SVM Performance Evaluation: Hold-out Method

```
predict_svm_linear <- predict(model_svm_linear, newdata = diabetes_test)
confusionMatrix(predict_svm_linear, diabetes_test$diabetes)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction neg pos
```

```
##           neg 134  36
```

```
##           pos  16  44
```

```
##
```

```
##           Accuracy : 0.7739
```

```
##           95% CI : (0.7143, 0.8263)
```

```
## No Information Rate : 0.6522
```

```
## P-Value [Acc > NIR] : 4.208e-05
```

```
##
```

```
##           Kappa : 0.4708
```

```
## McNemar's Test P-Value : 0.008418
```

```
##
```

```
##           Sensitivity : 0.8933
```

```
##           Specificity : 0.5500
```

```
##           Pos Pred Value : 0.7882
```

21/35

# SVM with Non-linear Kernel: RBF

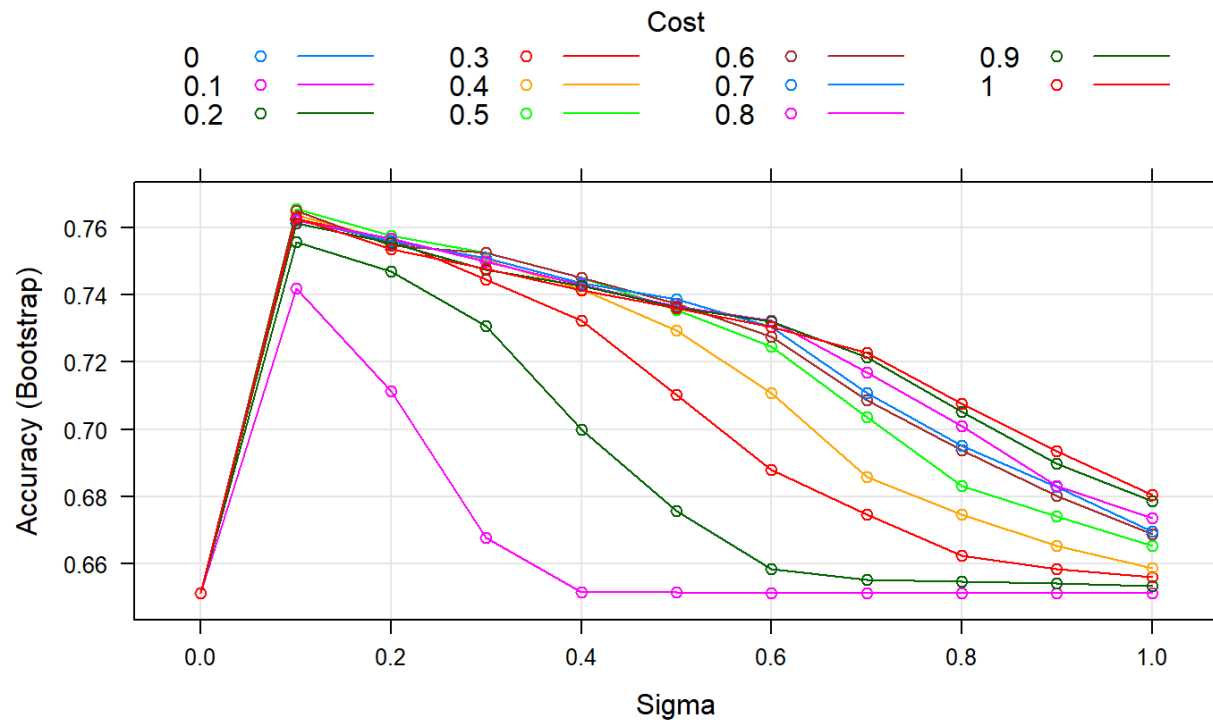
```
model_svm_rbf <- train(diabetes ~ ., data = diabetes_train,  
  preProcess = c("center", "scale"),  
  tuneGrid = expand.grid(sigma = seq(0, 1, 0.1),  
    C = seq(0, 1, 0.1)),  
  method = "svmRadial",  
  trControl = trainControl(method = "boot",  
    number = 25))
```

```
model_svm_rbf
```

```
## Support Vector Machines with Radial Basis Function Kernel
##
## 538 samples
## 8 predictor
## 2 classes: 'neg', 'pos'
##
## Pre-processing: centered (8), scaled (8)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 538, 538, 538, 538, 538, 538, ...
## Resampling results across tuning parameters:
##
##  sigma  C    Accuracy  Kappa
##  0.0    0.0      NaN      NaN
##  0.0    0.1  0.6512478  0.000000000
##  0.0    0.2  0.6512478  0.000000000
##  0.0    0.3  0.6512478  0.000000000
##  0.0    0.4  0.6512478  0.000000000
##  0.0    0.5  0.6512478  0.000000000
##  0.0    0.6  0.6512478  0.000000000
##  0.0    0.7  0.6512478  0.000000000
##  0.0    0.8  0.6512478  0.000000000
##  0.0    0.9  0.6512478  0.000000000
##  0.0    1.0  0.6512478  0.000000000
```

23/35

```
plot(model_svm_rbf)
```





```
predict_svm_rbf <- predict(model_svm_rbf, newdata = diabetes_test)
confusionMatrix(predict_svm_rbf, diabetes_test$diabetes)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction neg pos
```

```
##      neg 131  43
```

```
##      pos   19  37
```

```
##
```

```
##              Accuracy : 0.7304
```

```
##              95% CI : (0.6682, 0.7866)
```

```
## No Information Rate : 0.6522
```

```
## P-Value [Acc > NIR] : 0.006878
```

```
##
```

```
##              Kappa : 0.3611
```

```
## McNemar's Test P-Value : 0.003489
```

```
##
```

```
##              Sensitivity : 0.8733
```

```
##              Specificity : 0.4625
```

```
##              Pos Pred Value : 0.7529
```

```
##              Neg Pred Value : 0.6607
```

```
##              Prevalence : 0.6522
```

```
##              Detection Rate : 0.5696
```

```
##              Detection Prevalence : 0.7565
```

25/35

# Compare the Performance of Multiple Algorithms

```

model_comparison <- resamples(list(SVMLinear = model_svm_linear,
                                   SVMRBF = model_svm_rbf))

summary(model_comparison)

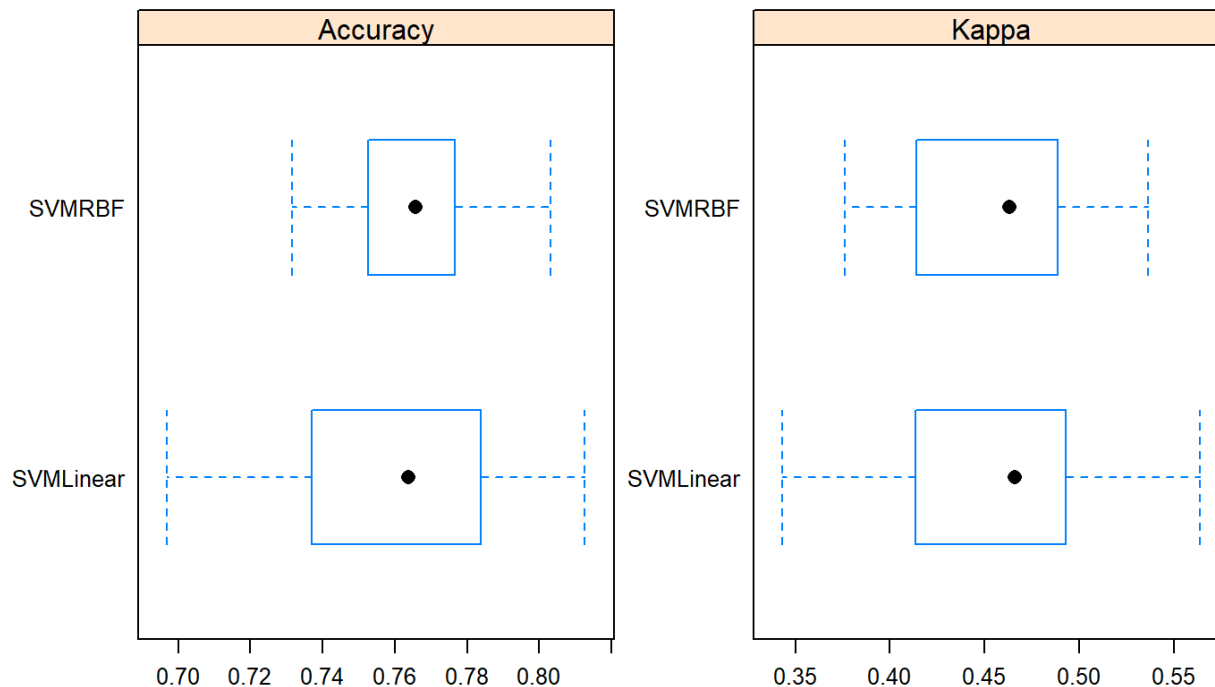
##
## Call:
## summary.resamples(object = model_comparison)
##
## Models: SVMLinear, SVMRBF
## Number of resamples: 25
##
## Accuracy
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## SVMLinear 0.6969697 0.7371134 0.7638191 0.7618265 0.7839196 0.8125000    0
## SVMRBF    0.7315789 0.7526316 0.7656250 0.7654968 0.7766497 0.8031088    0
##
## Kappa
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## SVMLinear 0.3432836 0.4134789 0.4659401 0.4582616 0.4931610 0.5637772    0
## SVMRBF    0.3760464 0.4141189 0.4633530 0.4528180 0.4885286 0.5365331    0

```

# Graphically Compare Performance

```
scales <- list(x = list(relation = "free"),  
               y = list(relation = "free"))
```

```
bwplot(model_comparison, scales = scales)
```



# Load in Libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import metrics
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import ShuffleSplit
from sklearn.model_selection import cross_val_score
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.svm import SVC
from sklearn import svm
from sklearn import model_selection
from sklearn.model_selection import GridSearchCV
from matplotlib.legend_handler import HandlerLine2D
```

# Data Splitting and Preparation

```
diabetes = r.diabetes
# diabetes = pd.read_csv("https://datahub.io/machine-learning/diabetes/r/diabetes.csv")
X = diabetes.drop("class", axis=1)
y = diabetes["class"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)
scaler = StandardScaler()
scaler.fit(X_train)

## StandardScaler(copy=True, with_mean=True, with_std=True)

X_train_std = scaler.transform(X_train)
X_test_std = scaler.transform(X_test)
```

# Linear SVM Model Tuning

```
lin_clf = svm.LinearSVC()  
lin_clf
```

```
## LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,  
##          intercept_scaling=1, loss='squared_hinge', max_iter=1000,  
##          multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,  
##          verbose=0)
```

```
bs = model_selection.ShuffleSplit(n_splits=25, test_size=0.3, random_state=0)  
param_grid = {'C':[0.25,0.5,0.75,1], 'penalty':['l2']}
```

```
gridbs = GridSearchCV(lin_clf, param_grid, cv=bs)  
gridbs.fit(X_train_std, y_train)
```

```
## GridSearchCV(cv=ShuffleSplit(n_splits=25, random_state=0, test_size=0.3, train_size=None),  
##            error_score='raise-deprecating',  
##            estimator=LinearSVC(C=1.0, class_weight=None, dual=True,  
##                               fit_intercept=True, intercept_scaling=1,  
##                               loss='squared_hinge', max_iter=1000,  
##                               multi_class='ovr', penalty='l2',  
##                               random_state=None, tol=0.0001, verbose=0),  
##            iid='warn', n_jobs=None,
```

30/35

# Hyper-Parameter $C$ Sensitivity

```
C = np.arange(0.05,1,0.05)
train_results_svm = []
test_results_svm = []
for n in C:
    model = svm.LinearSVC(C=n)
    model.fit(X_train_std, y_train)
    acc = cross_val_score(model, X_train_std, y_train, cv=bs).mean()*100
    train_results_svm.append(acc)

## LinearSVC(C=0.05, class_weight=None, dual=True, fit_intercept=True,
##           intercept_scaling=1, loss='squared_hinge', max_iter=1000,
##           multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
##           verbose=0)
## LinearSVC(C=0.1, class_weight=None, dual=True, fit_intercept=True,
##           intercept_scaling=1, loss='squared_hinge', max_iter=1000,
##           multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
##           verbose=0)
## LinearSVC(C=0.15000000000000002, class_weight=None, dual=True,
##           fit_intercept=True, intercept_scaling=1, loss='squared_hinge',
##           max_iter=1000, multi_class='ovr', penalty='l2', random_state=None,
##           tol=0.0001, verbose=0)
## LinearSVC(C=0.2, class_weight=None, dual=True, fit_intercept=True,
##           intercept_scaling=1, loss='squared_hinge', max_iter=1000,
```

31/35

# Linear SVM Performance Evaluation

```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import classification_report
y_pred = model.predict(X_test_std)
cmnb=confusion_matrix(y_test, y_pred, labels=["tested_negative","tested_positive"])
target_names = ["tested_negative","tested_positive"]
print(classification_report(y_test, y_pred, target_names=target_names))

```

```

##                precision    recall  f1-score   support
##
## tested_negative      0.79      0.90      0.84       146
## tested_positive      0.78      0.58      0.66       85
##
##      accuracy          0.78          231
##      macro avg      0.78      0.74      0.75       231
##      weighted avg      0.78      0.78      0.78       231

```



# SVM with RBF Kernel Function: $\gamma$

```
rbf_clf = svm.SVC()
gamma = [0.001,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1]
train_results_svm=[]
train_results_svmC2 = []
train_results_svmC3= []
for n in gamma:
    model1 = SVC(C=0.1,kernel='rbf',gamma=n,cache_size=2000)
    model1.fit(X_train,y_train)
    acc = cross_val_score(model1, X_train_std, y_train, cv=bs).mean()*100
    train_results_svm.append(acc)
    model2 = SVC(C=0.5,kernel='rbf',gamma=n,cache_size=2000)
    acc2 = cross_val_score(model2, X_train_std, y_train, cv=bs).mean()*100
    train_results_svmC2.append(acc2)
    model3 = SVC(C=1,kernel='rbf',gamma=n,cache_size=2000)
    acc3 = cross_val_score(model3, X_train_std, y_train, cv=bs).mean()*100
    train_results_svmC3.append(acc3)

## SVC(C=0.1, cache_size=2000, class_weight=None, coef0=0.0,
##      decision_function_shape='ovr', degree=3, gamma=0.001, kernel='rbf',
##      max_iter=-1, probability=False, random_state=None, shrinking=True,
##      tol=0.001, verbose=False)
## SVC(C=0.1, cache_size=2000, class_weight=None, coef0=0.0,
##      decision_function_shape='ovr', degree=3, gamma=0.1, kernel='rbf',
```

33/35

# RBM Kernel SVM Performance

```
param_grid = {'C':[0.5], 'gamma':[0.1]}
gridrbf = GridSearchCV(model2, param_grid, cv=bs)
gridrbf.fit(X_train_std, y_train)

## GridSearchCV(cv=ShuffleSplit(n_splits=25, random_state=0, test_size=0.3, train_size=None),
##             error_score='raise-deprecating',
##             estimator=SVC(C=0.5, cache_size=2000, class_weight=None, coef0=0.0,
##                           decision_function_shape='ovr', degree=3, gamma=1,
##                           kernel='rbf', max_iter=-1, probability=False,
##                           random_state=None, shrinking=True, tol=0.001,
##                           verbose=False),
##             iid='warn', n_jobs=None, param_grid={'C': [0.5], 'gamma': [0.1]},
##             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
##             scoring=None, verbose=0)
```

# RBM Kernel SVM Performance

```
y_pred2 = gridrbf.predict(X_test_std)
cmnb2=confusion_matrix(y_test, y_pred2, labels=["tested_negative", "tested_positive"])
target_names = ["tested_negative", "tested_positive"]
print(classification_report(y_test, y_pred2, target_names=target_names))
```

```
##                precision    recall  f1-score   support
##
## tested_negative      0.78      0.92      0.85       146
## tested_positive      0.81      0.55      0.66        85
##
##      accuracy              0.79       231
##      macro avg      0.80      0.74      0.75       231
##      weighted avg      0.79      0.79      0.78       231
```