

# Week 10 Statistical Data Mining

Theory and Practice

Ying Lin, Ph.D

March 27, 2020

# Overview of the Topics

- Linear regression
- Logistic regression and activation function
- Model coefficient estimation and interpretation
- Loss function, gradient descent, optimization
- R/Python demo

# Regression Motivation

- Linear Regression: model numerical outcome
  - Simple vs. multiple regression

$$\begin{aligned}f(x) &= \beta_0 + \beta_1 \cdot x_1 + \dots + \beta_i \cdot x_i + \dots + \beta_n \cdot x_n \\&= \sum_{i=1}^n \beta_i x_i \\&= \boldsymbol{\beta}^T \mathbf{x}\end{aligned}$$

- Logistic regression: Generalized linear regression
  - Model binary variable with two possible outcomes
  - Model multinomial categorical variables
- Key: estimation of model parameters which are coefficients in the regression model
  - Interpretation of coefficients

# Key Assumptions of Linear Regression

- Linear relationship between IVs and DVs
  - Checked by scatterplots for linear or curvilinear relationship
  - Transformation techniques such as normalizing DV
- Homoscedasticity: similar variance of error terms across space of IVs
  - Checked by scatterplot of residuals vs. predicted values
- Multivariate normality: residues are normally distributed
- Residual/error terms should be independent: no multicollinearity
  - Tested with Variance Inflation Factor (VIF) values

# R-Squared

- Percent of variance that could be explained by the model

$$TSS = \sum_i (y_i - \bar{y})^2$$

$$SSE = \sum_i (f(x_i) - y_i)^2$$

$$SSM = TSS - SSE$$

$$R^2 = \frac{SSM}{TSS}$$

# Ordinary Least Square (OLS)

- Minimize sum of squared error between the data points and the regression line/model
- How to estimate regression coefficients
  - Solve the model parameters analytically (closed-form solution)
  - Use the iterative optimization algorithms (Gradient Descent, Stochastic Gradient Descent, Newton's method, etc.)

# Ordinary Least Square (OLS)

- How to estimate the model coefficients (in specifications) that minimizes such squared difference between real value and predicted values

$$\frac{\partial S}{\partial a} = \frac{\partial \left( \sum (y - ax - b)^2 \right)}{\partial a} = 2 \sum (y - ax - b) \cdot (-x - 0)$$

the parameters of regression model  $\hat{y} = ax + b$  are  $a = \frac{n\bar{x}\bar{y} - \sum xy}{n\bar{x}^2 - \sum x^2}$  and  $b = \bar{y} - a\bar{x}$

# Objective Function

- Objective function: machine learning algorithms rely on minimizing or maximizing objective functions
- Loss function: the group of objective functions that are minimized; a measure of how good a prediction model does in terms of predicting the expected outcome
  - Regression loss
  - Classification loss



# Regression Loss Function

- Mean squared error (quadratic loss, L2 loss)

$$MSE = \frac{\sum_{i=1}^n (y_i - y_i^p)^2}{n}$$

- Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{MSE}$$

- Mean absolute error (L1 loss)

$$MAE = \frac{\sum_{i=1}^n |y_i - y_i^p|}{n}$$

# Comparison between L1 and L2 Loss Function

- MSE/RMSE is more sensitive towards outliers than MAE
  - If outliers represent anomaly that are important, then use MSE/RMSE
  - If outliers are due to erroneous data, then use MAE
  - L2 loss function works well in most of situations and consider removing outliers first
- MAE has larger and constant gradients unlike the dynamic gradient of MSE/RMSE which adjusts accordingly to the size of error
  - MAE requires an adaptive learning rate to complement

# Other Loss Functions

- Huber loss (Smooth Mean Absolute Error)


$$L_{\delta}(y, f(x)) = \begin{cases} \frac{1}{2}(y - f(x))^2 & \text{for } |y - f(x)| \leq \delta \\ \delta|y - f(x)| - \frac{1}{2}\delta^2 & \text{otherwise} \end{cases}$$

- Other considerations
  - Choice of hyperparameter  $\delta$  determines outliers: residues larger than  $\delta$  are minimized with L1; residues smaller than  $\delta$  are minimized with L2
  - Huber loss curves around the minima which decreases the gradient; and it is more robust to outliers than MSE

# Gradient Descent

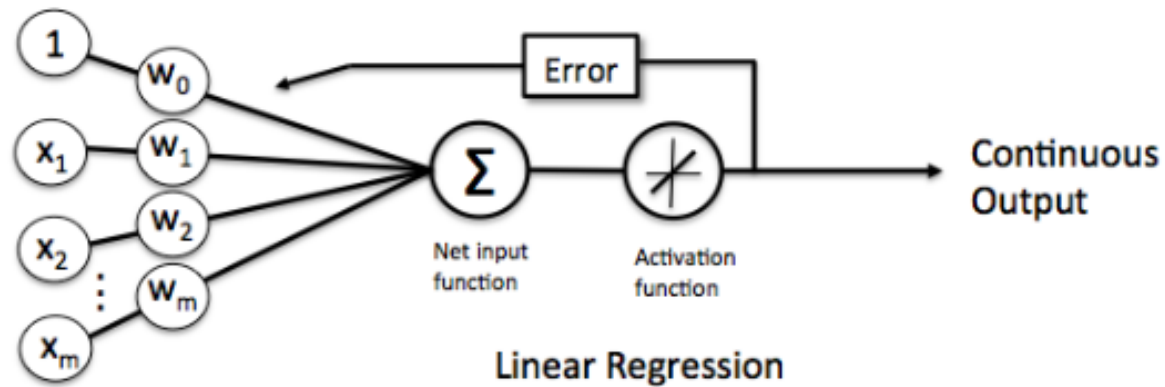
- Popular optimization algorithm to find the local minimum for a differentiable function, e.g. loss function
- An iterative process of minimizing a function by following the gradient (slope) of the cost function

$$m_{n+1} = m_n - \alpha \frac{\partial}{\partial m_n} LF_{m_n}$$



# Regression Model

- Linear regression



- How is a logistic regression different?
  - Activation function
  - Loss function

# Logistic/Sigmoid Function

- Logistic function

$$f(x) = \frac{1}{1 + e^{-k(x-x_0)}}$$

- Sigmoid function: a special case of logistic function or standard logistic function where  $K = 1$  and  $x_0 = 0$

$$S(x) = \frac{1}{1 + e^{-x}}$$

- Logit function: inverse of sigmoid function

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

$$\text{odd} = \frac{p}{1-p}$$

# Interpretation of Logistic Regression Coefficients

$$\text{logit}(p) = \sum_i^n \beta_i \cdot x_i + \beta_0$$

$$p = \frac{1}{1 + e^{-\sum_i^n \beta_i \cdot x_i + \beta_0}}$$


- How to interpret coefficient  $\beta_i$ ?

$$\log(\text{odd}_1) = \beta_0 + \beta_1 \cdot x_1 + \dots + \beta_i \cdot x_i + \dots + \beta_n \cdot x_n$$

$$\log(\text{odd}_2) = \beta_0 + \beta_1 \cdot x_1 + \dots + \beta_i \cdot (x_i + 1) + \dots + \beta_n \cdot x_n$$

$$\log(\text{odd}_2) - \log(\text{odd}_1) = \log \frac{\text{odd}_2}{\text{odd}_1} = \beta_i$$

# Sigmoid Function: S-Shape and its Properties





# Classification Loss Function

- Cross entropy loss function: negative log likelihood
  - $CrossEntropyLoss(y_i, \hat{y}_i) = -(y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$
  - Increases as the predicted probability diverges from the actual label
  - Property: penalize heavily the predictions that are  $\hat{y}_i = 0$  but  $y_i = 1$  or  $\hat{y}_i = 1$  but  $y_i = 0$
- Gradient of cross entropy loss (or log loss)
  - $\partial E / \partial p_i = -c_i / p_i + (1 - c_i) / (1 - p_i)$

# Regression with Regularization: Lasso vs Ridge

- Intuition: add a penalty term to loss function that is based on magnitude of coefficients as a way to reduce model complexity and overfitting, i.e. regularization
  - $RLS = LS + \lambda \sum_{i=1}^n \beta_i^q$
- Lasso regularization: L1 loss function ( $q = 1$ )
  - Dimension reduction and feature selection
- Ridge regularization: L2 loss function ( $q = 2$ )
  - Feature ranking and relative importance; but can't zero out coefficients like Lasso
- Shrinkage parameter:  $\lambda \geq 0$ 
  - $\lambda = 0$ : no regularization, just regular regression
  - As  $\lambda$  increases, the coefficients decrease

# Estimation of Coefficients of Logistic Regression

- Statistical approach: Maximum Likelihood Estimation (MLE)
  - Estimate the distribution parameters ( $\theta$ ) that maximize the product of probabilities of observing all the data points
- Machine learning approach
  - Loss function: cross-entropy function
  - Optimization algorithm, e.g. stochastic gradient descent (SGD)
- Difference between two approaches

# Demo Data: Boston Housing Price

```
library(caret)
library(mlbench)
data("BostonHousing")
str(BostonHousing)
```

```
## 'data.frame':    506 obs. of  14 variables:
## $ crim      : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn        : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus     : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas      : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ nox       : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm        : num  6.58 6.42 7.18 7 7.15 ...
## $ age       : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis       : num  4.09 4.97 4.97 6.06 6.06 ...
## $ rad       : num  1 2 2 3 3 3 5 5 5 5 ...
## $ tax       : num  296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio   : num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ b         : num  397 397 393 395 397 ...
## $ lstat     : num  4.98 9.14 4.03 2.94 5.33 ...
## $ medv      : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

# Train Linear Model

```
model_lm <- train(medv ~ ., data = BostonHousing, method = "lm")
print(model_lm)

## Linear Regression
##
## 506 samples
## 13 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 506, 506, 506, 506, 506, 506, ...
## Resampling results:
##
## RMSE      Rsquared    MAE
## 4.731782  0.7207107  3.378688
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

# Model Specification

```
model_lm$finalModel
```

```
##
```

```
## Call:
```

```
## lm(formula = .outcome ~ ., data = dat)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      crim          zn          indus          chas1
##  3.646e+01   -1.080e-01   4.642e-02   2.056e-02   2.687e+00
##          nox          rm          age          dis          rad
## -1.777e+01   3.810e+00   6.922e-04  -1.476e+00   3.060e-01
##          tax      ptratio          b          lstat
## -1.233e-02  -9.527e-01   9.312e-03  -5.248e-01
```

# Examine the Linear Regression Model

```
names(summary(model_lm))
```

```
## [1] "call"          "terms"          "residuals"      "coefficients"
## [5] "aliased"        "sigma"          "df"             "r.squared"
## [9] "adj.r.squared" "fstatistic"     "cov.unscaled"
```

```
summary(model_lm)
```

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-15.595	-2.730	-0.518	1.777	26.199

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	3.646e+01	5.103e+00	7.144	3.28e-12	***
crim	-1.080e-01	3.286e-02	-3.287	0.001087	**
zn	4.642e-02	1.373e-02	3.382	0.000778	***
indus	2.056e-02	6.150e-02	0.334	0.738288	

23/41

# Examine the Linear Regression Model (2)

```
coef <- summary(model_lm)$coefficients
coef <- coef[coef[, 4] < 0.05, ]
coef[order(coef[, 4]), ]
```

##		Estimate	Std. Error	t value	Pr(> t )
##	lstat	-0.524758378	0.050715278	-10.347146	7.776912e-23
##	rm	3.809865207	0.417925254	9.116140	1.979441e-18
##	dis	-1.475566846	0.199454735	-7.398004	6.013491e-13
##	ptratio	-0.952747232	0.130826756	-7.282511	1.308835e-12
##	(Intercept)	36.459488385	5.103458811	7.144074	3.283438e-12
##	nox	-17.766611228	3.819743707	-4.651257	4.245644e-06
##	rad	0.306049479	0.066346440	4.612900	5.070529e-06
##	b	0.009311683	0.002685965	3.466793	5.728592e-04
##	zn	0.046420458	0.013727462	3.381576	7.781097e-04
##	crim	-0.108011358	0.032864994	-3.286517	1.086810e-03
##	tax	-0.012334594	0.003760536	-3.280009	1.111637e-03
##	chas1	2.686733819	0.861579756	3.118381	1.925030e-03



# Information of the Attributes

- MEDV - Median value of owner-occupied homes in \$1000's
- CRIM - per capita crime rate by town
- ZN - proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS - proportion of non-retail business acres per town.
- CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise)
- NOX - nitric oxides concentration (parts per 10 million)
- RM - average number of rooms per dwelling
- AGE - proportion of owner-occupied units built prior to 1940
- DIS - weighted distances to five Boston employment centres
- RAD - index of accessibility to radial highways
- TAX - full-value property-tax rate per \$10,000
- PTRATIO - pupil-teacher ratio by town
- B -  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of blacks by town
- LSTAT - % lower status of the population

# Model Details

```
summary(model_lm)
```

```
##
```

```
## Call:
```

```
## lm(formula = .outcome ~ ., data = dat)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -15.595  -2.730  -0.518   1.777  26.199
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.646e+01  5.103e+00   7.144 3.28e-12 ***
## crim        -1.080e-01  3.286e-02  -3.287 0.001087 **
## zn           4.642e-02  1.373e-02   3.382 0.000778 ***
## indus        2.056e-02  6.150e-02   0.334 0.738288
## chas1        2.687e+00  8.616e-01   3.118 0.001925 **
## nox          -1.777e+01  3.820e+00  -4.651 4.25e-06 ***
## rm           3.810e+00  4.179e-01   9.116 < 2e-16 ***
## age          6.922e-04  1.321e-02   0.052 0.958229
## dis          -1.476e+00  1.995e-01  -7.398 6.01e-13 ***
## rad           3.060e-01  6.635e-02   4.613 5.07e-06 ***
## tax          -1.233e-02  3.760e-03  -3.280 0.001112 **
```

26/41

# Model Residuals' Plot

```
plot(model_lm$finalModel)
```

# Run LM Model with Standardized Data

```
model_lm2 <- train(medv ~ ., data = BostonHousing, method = "lm", preProcess = c("center", "scale"),  
print(model_lm2)
```

```
## Linear Regression  
##  
## 506 samples  
## 13 predictor  
##  
## Pre-processing: centered (13), scaled (13)  
## Resampling: Bootstrapped (25 reps)  
## Summary of sample sizes: 506, 506, 506, 506, 506, 506, ...  
## Resampling results:  
##  
## RMSE      Rsquared    MAE  
## 4.913235  0.7197477  3.450471  
##  
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

# Calculate R Square and Compare with Output (1)

```
predict_lm <- predict(model_lm, newdata = BostonHousing, type = "raw")
real_predict <- data.frame(cbind(real = BostonHousing$medv, predict = predict_lm,
                                square_diff = (BostonHousing$medv - predict_lm)^2))
str(real_predict)
```

```
## 'data.frame':    506 obs. of  3 variables:
## $ real          : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
## $ predict       : num  30 25 30.6 28.6 27.9 ...
## $ square_diff: num  36 11.7 17.1 23 68.2 ...
```

# Calculate R Square and Compare with Output (2)

```
SSE <- sum(real_predict$square_diff)
TSS <- var(BostonHousing$medv) * (length(BostonHousing$medv) + 1)
R_squared <- 1 - SSE/TSS
R_squared
```

```
## [1] 0.7416658
```

```
summary(model_lm)$r.squared
```

```
## [1] 0.7406427
```

# Logistic Regression

```
library(arules)
BostonHousing2 <- BostonHousing
BostonHousing2$medv <- discretize(BostonHousing2$medv,
                                   method = "frequency",
                                   breaks = 2,
                                   labels = c("low", "high"))

table(BostonHousing2$medv)

##
##  low high
##  251  255

model_glm <- train(medv ~ ., data = BostonHousing2,
                   method = "glm", family = "binomial")

print(model_glm)

## Generalized Linear Model
##
## 506 samples
## 13 predictor
## 2 classes: 'low', 'high'
```

31/41

# Logistic Regression Model Output

```
summary(model_glm)
```

```
##
```

```
## Call:
```

```
## NULL
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -2.0015  -0.3574   0.0085   0.2981   3.3286
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) 12.065836   4.007916   3.011  0.00261 **
## crim        -0.082159   0.074950  -1.096  0.27300
## zn           0.012275   0.013525   0.908  0.36411
## indus        0.029454   0.043259   0.681  0.49594
## chas1        1.659713   0.664634   2.497  0.01252 *
## nox          -7.283784   2.733645  -2.664  0.00771 **
## rm           1.617226   0.440282   3.673  0.00024 ***
## age         -0.028603   0.010451  -2.737  0.00620 **
## dis          -0.711035   0.168061  -4.231 2.33e-05 ***
## rad           0.237470   0.060920   3.898 9.70e-05 ***
## tax          -0.008461   0.002919  -2.899  0.00374 **
```

32/41



# Relative Importance of Variables by Logistic Regression

```
varImp(model_glm)
```

```
## glm variable importance
```

```
##
```

```
##          Overall
```

```
## lstat    100.000
```

```
## ptratio  86.431
```

```
## dis      68.821
```

```
## rad      62.370
```

```
## rm       58.010
```

```
## tax      43.000
```

```
## age      39.859
```

```
## nox      38.455
```

```
## chas1    35.212
```

```
## b        17.748
```

```
## crim     8.051
```

```
## zn       4.394
```

```
## indus    0.000
```

# Import Python Libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import load_boston
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.pipeline import make_pipeline
```

# Data Loading and Preprocessing

```
boston = load_boston()
boston_data = pd.DataFrame(boston.data, columns = boston.feature_names)
boston_data['Price'] = boston.target
X = boston_data.drop('Price', axis=1)
y = boston_data['Price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=16)
boston_data.info()
```

```
## <class 'pandas.core.frame.DataFrame'>
## RangeIndex: 506 entries, 0 to 505
## Data columns (total 14 columns):
## CRIM      506 non-null float64
## ZN        506 non-null float64
## INDUS     506 non-null float64
## CHAS      506 non-null float64
## NOX       506 non-null float64
## RM        506 non-null float64
## AGE       506 non-null float64
## DIS       506 non-null float64
## RAD       506 non-null float64
## TAX       506 non-null float64
## PTRATIO   506 non-null float64
## B         506 non-null float64
```

35/41

# Linear Regression Training and Testing


```
lr_pipe = LinearRegression()  
lr_pipe.fit(X_train, y_train)  
  
## LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)  
  
y_pred = lr_pipe.predict(X_test)  
print(f"RMSE: {round(np.sqrt(metrics.mean_squared_error(y_test, y_pred)), 3)}")  
  
## RMSE: 4.614
```

# Visualize Regression Line

```
plt.figure(figsize=(4, 3))  
plt.scatter(y_test, y_pred)  
plt.plot([0, 50], [0, 50], '--k')  
plt.axis('tight')
```

```
## (-2.500483870967742, 52.51016129032258, -10.282281648953647, 52.87058484042636)
```

```
plt.tight_layout()  
plt.xlabel('True price ($1000s)')  
plt.ylabel('Predicted price ($1000s)')  
plt.show()
```



# Linear Model Coefficient Estimation

```
for idx, col_name in enumerate(X_train.columns):  
    print(f"The coefficient for {col_name} is {lr_pipe.coef_[idx]}")
```

```
## The coefficient for CRIM is -0.135721389191917  
## The coefficient for ZN is 0.03531958631459615  
## The coefficient for INDUS is 0.0020350011834421146  
## The coefficient for CHAS is 2.425658595197517  
## The coefficient for NOX is -15.996609853279406  
## The coefficient for RM is 3.997523842035774  
## The coefficient for AGE is 0.022861021101290676  
## The coefficient for DIS is -1.3542588201878618  
## The coefficient for RAD is 0.346279075866271  
## The coefficient for TAX is -0.011431344325432802  
## The coefficient for PTRATIO is -1.0161040881448886  
## The coefficient for B is 0.010906415198789543  
## The coefficient for LSTAT is -0.6351622221570083
```

```
plt.scatter(lr_pipe.predict(X_train), lr_pipe.predict(X_train)-y_train, c='b', s=40, alpha=0.5)  
plt.hlines(y=0, xmin=0, xmax=50)  
plt.xlabel("Fitted")  
plt.ylabel("Residuals")  
plt.title("Residuals vs. fitted")  
plt.show()
```

38/41

# Evaluate Linear Regression Model

```
data_tuples = list(zip(y_test, y_pred))
real_predict = pd.DataFrame(data_tuples, columns=['Real', 'Predict'])
real_predict['squared_dif'] = (real_predict["Real"]-real_predict["Predict"])**2
real_predict.head(3)
```

```
##      Real      Predict  squared_dif
## 0  23.3  27.406211    16.860968
## 1  12.8  13.179976     0.144382
## 2   6.3   9.676904    11.403484
```

```
SSE = sum(real_predict['squared_dif'])
TSS = np.var(y_test) * (len(y_test)+1)
real_predict.head(3)
```

```
##      Real      Predict  squared_dif
## 0  23.3  27.406211    16.860968
## 1  12.8  13.179976     0.144382
## 2   6.3   9.676904    11.403484
```

$R_{\text{squared}} = 1 - \text{SSE}/\text{TSS}$

$R_{\text{squared}}$

# Prepare for Logistic Regression Dataset

```
boston_data2 = boston_data.copy()
boston_data2['Price'] = pd.qcut(boston_data2['Price'], 2, labels=["low", "high"])
X2 = boston_data2.drop('Price', axis=1)
y2 = boston_data2['Price']
X_train2, X_test2, y_train2, y_test2 = train_test_split(X2, y2, test_size=0.3, random_state=16)
boston_data2.Price.value_counts()

## low      256
## high     250
## Name: Price, dtype: int64
```



# Logistic Regression Model Training and Testing

```
logr_pipe = make_pipeline(StandardScaler(), LogisticRegression(solver='lbfgs'))  
logr_pipe.fit(X_train2, y_train2)
```

```
## Pipeline(memory=None,  
##         steps=[('standardscaler',  
##               StandardScaler(copy=True, with_mean=True, with_std=True)),  
##               ('logisticregression',  
##               LogisticRegression(C=1.0, class_weight=None, dual=False,  
##                               fit_intercept=True, intercept_scaling=1,  
##                               l1_ratio=None, max_iter=100,  
##                               multi_class='warn', n_jobs=None,  
##                               penalty='l2', random_state=None,  
##                               solver='lbfgs', tol=0.0001, verbose=0,  
##                               warm_start=False))],  
##         verbose=False)
```

```
y_pred2 = logr_pipe.predict(X_test2)  
y_pred2
```

```
## array(['high', 'low', 'low', 'high', 'low', 'high', 'high', 'low', 'low',  
##       'high', 'low', 'low', 'low', 'high', 'low', 'high', 'low', 'high',  
##       'low', 'low', 'low', 'low', 'low', 'low', 'low', 'low', 'low',
```

41/41